

AVaaS : Agent Validator

Coforge TechCon 2026 ● ●

Vision Statement: AVaaS is a multi-tenant platform to test, evaluate, compare, and monitor AI agents before release and in production-like environments. It combines automated test runs, trace observability, regression gates, and alerting into one workflow. The goal is to help teams ship agentic AI faster with confidence, reliability, and governance.



Industry : QE | Sub-Industry : QA Automation

Title : Agent Validator Platform

The Problem: Silent Failures

Agents hallucinate API arguments, attempt unapproved tool calls, or output malformed JSON payloads.

The Problem: Regression

Modifying a system prompt, RAG context, or model version silently degrades agent performance on edge cases.

1-Liner Use Cases

1. **Real-Time Testing:** Execute multi-turn agent tests with trace interception.
2. **Dual-Tier Scoring:** Grade responses using rules + LLM-as-a-Judge.
3. **Regression Gates:** Compare baselines to auto-detect degradation.
4. **Synthetic Generation:** Auto-generate adversarial edge cases.

MVP Scope & Tech Stack



Hackathon MVP Scope

Ingestion: Accept any FastAPI Agent Endpoint.

Execution: High-concurrency async load testing.

Governance: Baseline vs. Candidate builds to trigger Pass/Fail release gates.

- UI: React
- Backend: FastAPI
- Eval: DeepEval
- Trace & Metrics: Langfuse & OpenTelemetry / LangSmith
- LLM: Ollama / Gemini

Architected as an internal tool for Coforge's QE practice, AVaaS uses a primary Open-Source stack (Ollama, DeepEval, Langfuse, OpenTelemetry) for client data sovereignty. It integrates commercial fallbacks (Gemini API, LangSmith) to guarantee zero-downtime execution and complete observability during agent evaluation.

Execution Architecture

User/QA Engineer

Web Dashboard, Agent Config
(End Point, System Prompt, Tools)

Phase 1 : Test Generation

Agent spec & Tool parser

LLM Test Generator

Test Cases(Normal, Edge,
Injection, Multi-Turn)

Phase 2 : Async Execution

Async Concurrent Runner

Target Agent Endpoint

Trace Collector(Collect traces
Tool calls, latency, Tokens)

Phase 3 : Evaluation (Dual-Tier)

Rule Based Judge

LLM as a Judge

Composite Score

Phase 4 : Regression & Reporting

Baseline Comparison

Failure analysis

Reports & Dashboard

Web Dashboard/Reports

View Results , Analytics

Relevance & Feasibility

Data Sovereignty: The open-source-first approach ensures client IP and system prompts never leave the corporate firewall.

Shift-Left Governance: AVaaS elevates Coforge's QE practice from legacy script maintenance to continuous agentic governance, providing automated release decisions before deployment.

Engineering the Autonomous Enterprise



Thank You



Industry : QE | Sub-Industry : QA Automation

Title : Agent Validator Platform